


## Setup Environment

```
# Install specific libraries
! pip install transformers
! pip install pycaret
```

 [Show hidden output](#)

```
import numpy as np
import pandas as pd
import pycaret
import transformers
from transformers import AutoModel, BertTokenizerFast
import matplotlib.pyplot as plt
from sklearn.metrics import plot_confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import torch
import torch.nn as nn
# specify GPU
device = torch.device("cuda")

# Mount Google Drive - applicable, if working on Google Drive
from google.colab import drive
drive.mount('/content/drive')
```

 Mounted at /content/drive

```
# Set Working Directory - if working on Google Drive
%cd /content/drive/MyDrive/1_LiveProjects/Project11_FakeNewsDetection
```

```
# # Set Working Directory - if working on Local Machine
# import os
# os.chdir('/Users//replace_me')
```

 /content/drive/MyDrive/Project11\_FakeNewsDetection

## Load Dataset

```
# Load Dataset
true_data = pd.read_csv('a1_True.csv')
fake_data = pd.read_csv('a2_Fake.csv')

# Generate labels True/Fake under new Target Column in 'true_data' and 'fake_data'
true_data['Target'] = ['True']*len(true_data)
fake_data['Target'] = ['Fake']*len(fake_data)

# Merge 'true_data' and 'fake_data', by random mixing into a single df called 'data'
data = true_data.append(fake_data).sample(frac=1).reset_index().drop(columns=['index'])


# See how the data looks like
print(data.shape)
data.head()
```

 (44898, 5)

		title	text	subject	date	Target
0	PROOF That Obama Interfered TWICE In Foreign E...	President Barack Obama recently said he told R...		left-news	Jun 8, 2017	Fake
1	Kremlin: no firm date yet for proposed congres...	MOSCOW (Reuters) - No firm date has been set y...		worldnews	November 28, 2017	True
2	TOBY KEITH Has AWESOME Response To Crybaby Att...	Country singer Toby Keith won t be bullied int...		left-news	Jan 16, 2017	Fake
3	BREAKING: Putin Tramples Obama's Imaginary Red...	We have a weak leader in our White House who s...		Government News	Sep 30, 2015	Fake


```
# Target column is made of string values True/Fake, let's change it to numbers 0/1 (Fake=1)
data['label'] = pd.get_dummies(data.Target)['Fake']
```

```
data.head()
```




	title	text	subject	date	Target	label
0	PROOF That Obama Interfered TWICE In Foreign E...	President Barack Obama recently said he told R...	left-news	Jun 8, 2017	Fake	1
1	Kremlin: no firm date yet for proposed congres...	MOSCOW (Reuters) - No firm date has been set y...	worldnews	November 28, 2017	True	0
2	TOBY KEITH Has AWESOME Response To Crybaby Att...	Country singer Toby Keith won t be bullied int...	left-news	Jan 16, 2017	Fake	1
3	BREAKING: Putin Tramples Obama's Imaginary Red...	We have a weak leader in our White House who s...	Government News	Sep 30, 2015	Fake	1
4	ULTIMATE HYPOCRITES! RUSSIAN Ambassador Visite...	According to the New Yorker, Bill Clinton rece...	politics	Mar 3, 2017	Fake	1

```
# Checking if our data is well balanced
label_size = [data['label'].sum(),len(data['label'])-data['label'].sum()]
plt.pie(label_size,explode=[0.1,0.1],colors=['firebrick','navy'],startangle=90,shadow=True,labels=['Fake','True'],autopct='%1.1f%%')
```



```
(<matplotlib.patches.Wedge at 0x7ff8f2d78810>,
 <matplotlib.patches.Wedge at 0x7ff8f2d8b310>],
 [Text(-1.1968727067385088, -0.0865778485782335, 'Fake'),
  Text(1.1968726986325005, 0.08657796063754254, 'True')],
 [Text(-0.6981757455974634, -0.05050374500396954, '52.3%'),
  Text(0.6981757408689586, 0.05050381037189981, '47.7%')])
```




> Train-test-split

[ ] 1 cell hidden

✓ BERT Fine-tuning

✓ Load pretrained BERT Model

```
# Load BERT model and tokenizer via HuggingFace Transformers
bert = AutoModel.from_pretrained('bert-base-uncased')
tokenizer = BertTokenizerFast.from_pretrained('bert-base-uncased')
```




```
Downloading config.json: 100% 570/570 [00:00<00:00, 15.7kB/s]
Downloading pytorch_model.bin: 100% 420M/420M [00:07<00:00, 61.1MB/s]
Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.seq_relationship.bias', 'cls.
- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initial
Downloading tokenizer_config.json: 100% 28.0/28.0 [00:00<00:00, 920B/s]
Downloading vocab.txt: 100% 226k/226k [00:00<00:00, 1.06MB/s]
Downloading tokenizer.json: 100% 455k/455k [00:00<00:00, 1.61MB/s]
```

✓ Prepare Input Data

```
# Plot histogram of the number of words in train data 'title'
seq_len = [len(title.split()) for title in train_text]

pd.Series(seq_len).hist(bins = 40,color='firebrick')
plt.xlabel('Number of Words')
plt.ylabel('Number of texts')
```

↗ Text(0, 0.5, 'Number of texts')



```
# BERT Tokenizer Functionality
sample_data = ["Build fake news model.",
               "Using bert."]
tokenized_sample_data = tokenizer.batch_encode_plus(sample_data,
                                                    padding=True)
print(tokenized_sample_data)
```

# Ref: <https://huggingface.co/docs/transformers/preprocessing>

↗ {'input\_ids': [[101, 3857, 8275, 2739, 2944, 1012, 102], [101, 2478, 14324, 1012, 102, 0, 0]], 'token\_type\_ids': [[0, 0, 0, 0, 0, 0, 0],

# Majority of titles above have word length under 15. So, we set max title length as 15

MAX LENGHT = 15

# Tokenize and encode sequences in the train set

```
tokens_train = tokenizer.batch_encode_plus(
    train_text.tolist(),
    max_length = MAX LENGHT,
    pad_to_max_length=True,
    truncation=True
)
```

# tokenize and encode sequences in the validation set

```
tokens_val = tokenizer.batch_encode_plus(
    val_text.tolist(),
    max_length = MAX LENGHT,
    pad_to_max_length=True,
    truncation=True
)
```

# tokenize and encode sequences in the test set

```
tokens_test = tokenizer.batch_encode_plus(
    test_text.tolist(),
    max_length = MAX LENGHT,
    pad_to_max_length=True,
    truncation=True
)
```

↗ /usr/local/lib/python3.7/dist-packages/transformers/tokenization\_utils\_base.py:2329: FutureWarning: The `pad\_to\_max\_length` argument is deprecated, will be removed in a future version.

# Convert lists to tensors

```
train_seq = torch.tensor(tokens_train['input_ids'])
train_mask = torch.tensor(tokens_train['attention_mask'])
train_y = torch.tensor(train_labels.tolist())
```

```
val_seq = torch.tensor(tokens_val['input_ids'])
val_mask = torch.tensor(tokens_val['attention_mask'])
val_y = torch.tensor(val_labels.tolist())
```

```
test_seq = torch.tensor(tokens_test['input_ids'])
test_mask = torch.tensor(tokens_test['attention_mask'])
test_y = torch.tensor(test_labels.tolist())
```

```
# Data Loader structure definition
from torch.utils.data import TensorDataset, DataLoader, RandomSampler, SequentialSampler
batch_size = 32 #define a batch size

train_data = TensorDataset(train_seq, train_mask, train_y) # wrap tensors
train_sampler = RandomSampler(train_data) # sampler for sampling the data during training
train_dataloader = DataLoader(train_data, sampler=train_sampler, batch_size=batch_size)
# dataLoader for train set

val_data = TensorDataset(val_seq, val_mask, val_y) # wrap tensors
val_sampler = SequentialSampler(val_data) # sampler for sampling the data during training
val_dataloader = DataLoader(val_data, sampler = val_sampler, batch_size=batch_size)
# dataLoader for validation set
```

## Freeze Layers


```
# Freezing the parameters and defining trainable BERT structure
for param in bert.parameters():
    param.requires_grad = False # false here means gradient need not be computed
```

## Define Model Architecture

```
class BERT_Arch(nn.Module):
    def __init__(self, bert):
        super(BERT_Arch, self).__init__()
        self.bert = bert
        self.dropout = nn.Dropout(0.1) # dropout layer
        self.relu = nn.ReLU() # relu activation function
        self.fc1 = nn.Linear(768,512) # dense layer 1
        self.fc2 = nn.Linear(512,2) # dense layer 2 (Output layer)
        self.softmax = nn.LogSoftmax(dim=1) # softmax activation function
    def forward(self, sent_id, mask):
        cls_hs = self.bert(sent_id, attention_mask=mask)['pooler_output']
        # pass the inputs to the model

        x = self.fc1(cls_hs)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.fc2(x) # output layer
        x = self.softmax(x) # apply softmax activation
        return x

model = BERT_Arch(bert)
# Defining the hyperparameters (optimizer, weights of the classes and the epochs)
# Define the optimizer
from transformers import AdamW
optimizer = AdamW(model.parameters(),
                    lr = 1e-5) # learning rate
# Define the loss function
cross_entropy = nn.NLLLoss()
# Number of training epochs
epochs = 2
```

 /usr/local/lib/python3.7/dist-packages/transformers/optimization.py:310: FutureWarning: This implementation of AdamW is deprecated and will be removed in a future version. Please use the implementation in torch.nn.optim.AdamW instead. Triggered by <transformers.optimization.AdamW object> at 0x7f8b1c1c1c1c.

## Define Train & Evaluate Function

```

# Defining training and evaluation functions
def train():
    model.train()
    total_loss, total_accuracy = 0, 0

    for step, batch in enumerate(train_dataloader):
        # iterate over batches
        if step % 50 == 0 and not step == 0:
            # progress update after every 50 batches.
            print(' Batch {:>5,} of {:>5,}'.format(step, len(train_dataloader)))
        batch = [r for r in batch]
        # push the batch to gpu
        sent_id, mask, labels = batch
        model.zero_grad()
        # clear previously calculated gradients
        preds = model(sent_id, mask)
        # get model predictions for current batch
        loss = cross_entropy(preds, labels)
        # compute loss between actual & predicted values
        total_loss = total_loss + loss.item()
        # add on to the total loss
        loss.backward()
        # backward pass to calculate the gradients
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
        # clip gradients to 1.0. It helps in preventing exploding gradient problem
        optimizer.step()
        # update parameters
        preds = preds.detach().cpu().numpy()
        # model predictions are stored on GPU. So, push it to CPU

    avg_loss = total_loss / len(train_dataloader)
    # compute training loss of the epoch
    # reshape predictions in form of (# samples, # classes)
    return avg_loss, preds

def evaluate():
    print("\nEvaluating...")
    model.eval()
    # Deactivate dropout layers
    total_loss, total_accuracy = 0, 0
    for step, batch in enumerate(val_dataloader):
        # Iterate over batches
        if step % 50 == 0 and not step == 0:
            # Progress update every 50 batches.
            # Calculate elapsed time in minutes.
            # Elapsed = format_time(time.time() - t0)
            print(' Batch {:>5,} of {:>5,}'.format(step, len(val_dataloader)))
            # Report progress
            batch = [t for t in batch]
            # Push the batch to GPU
            sent_id, mask, labels = batch
            with torch.no_grad():
                # Deactivate autograd
                preds = model(sent_id, mask)
                # Model predictions
                loss = cross_entropy(preds, labels)
                # Compute the validation loss between actual and predicted values
                total_loss = total_loss + loss.item()
                preds = preds.detach().cpu().numpy()
            avg_loss = total_loss / len(val_dataloader)
            # compute the validation loss of the epoch
    return avg_loss, preds

```

## ✓ Model training

```

# Train and predict
best_valid_loss = float('inf')
train_losses=[]
valid_losses=[]
# empty lists to store training and validation loss of each epoch

for epoch in range(epochs):
    print('\n Epoch {:} / {:}'.format(epoch + 1, epochs))
    train_loss = train()
    # train model
    valid_loss = evaluate()
    # evaluate model
    if valid_loss < best_valid_loss:
        # save the best model
        best_valid_loss = valid_loss
        torch.save(model.state_dict(), 'c2_new_model_weights.pt')
    train_losses.append(train_loss)
    # append training and validation loss
    valid_losses.append(valid_loss)

    print(f'\nTraining Loss: {train_loss:.3f}')
    print(f'\nValidation Loss: {valid_loss:.3f}')

```

## ✓ Model performance

```


# load weights of best model
path = 'c1_fakenews_weights.pt'
model.load_state_dict(torch.load(path))

```

→ <All keys matched successfully>

```
with torch.no_grad():
    preds = model(test_seq, test_mask)
    preds = preds.detach().cpu().numpy()

preds = np.argmax(preds, axis = 1)
print(classification_report(test_y, preds))
```



	precision	recall	f1-score	support
0	0.84	0.92	0.88	3213
1	0.92	0.84	0.88	3522
accuracy			0.88	6735
macro avg	0.88	0.88	0.88	6735
weighted avg	0.88	0.88	0.88	6735

> Fake News Predictions

[ ] ↳ 2 cells hidden